


YELLOW3 RESEARCH / MODEL INTELLIGENCE

THE MODEL INTELLIGENCE MONTHLY DECISION REPORT

Asia now captures 60.6% of observed routed AI traffic.

US share fell 2.73 percentage points. Europe remains below 1%.



Source: yellow3 Model Intelligence, live adoption dataset.
Data cut: 13 Jul 2026, 10:42 UTC. Shares shown as published in the July data snapshot.

ADOPTION / MODELS / PROVIDERS / ECONOMICS / POLICY AND ACCESS

Publication record.

The identity, methodology state and correction status attached to this edition.
These records remain with every archived version.

EDITION IDENTITY

PUBLICATION	PUBLICATION ID
Model Intelligence Monthly Decision Report	Y3-MI-2026-07-v1.0
PUBLISHER	REPORT STANDARD
yellow3 lab ApS	Y3-RS-v1.0
EDITION	METHODOLOGY
July 2026	MI-MET-v1.0
PUBLISHED	DATA SNAPSHOT
14 Jul 2026	MI-DATA-2026-07-13
REPORT LENGTH	DATA CUT
42 pages	13 Jul 2026, 10:42 UTC

EVIDENCE CLASSIFICATION

CONFIDENCE APPEARS ONLY ON ESTIMATED AND INFERRED MATERIAL

OBSERVED	Directly measured or recorded.
CALCULATED	Derived by yellow3 from disclosed inputs.
REPORTED	Claimed by a provider or institution.
ESTIMATED	Produced using a disclosed estimation method.
INFERRED	Editorial assessment based on multiple sources.

EDITION STATUS

COMPARABILITY	METHODOLOGY	CORRECTIONS
STANDARD SERIES	MI-MET-v1.0	NONE RECORDED
CITATION	CORRECTIONS	
yellow3 Model Intelligence Monthly Decision Report, Jul 2026, Y3-MI-2026-07-v1.0, page and figure number.	Corrections are never made silently. Revised editions receive a new version number and a permanent correction record.	

Contents.

Every edition uses the same 42-page structure.
Sections remain fixed so evidence can be compared page by page.

SECTION		PAGES
00	Front matter Cover, publication record, contents, reading guide and methodology status	01-04
01	Monthly Decision Briefing The decisions, signals and changes that matter this month	05-08
02	Adoption and Routed Usage Regional adoption, routed usage and movement over time	09-14
03	Model Landscape and Momentum Rankings, entrants, momentum and model-level change	15-20
04	Provider Strategy and Market Structure Provider position, concentration and strategic movement	21-25
05	Economics and Cost of Use Pricing, workload economics and comparable cost	26-31
06	Capability, Reliability and Enterprise Fit Capability, security, reliability and operational suitability	32-35
07	Policy, Access and Compute Sovereignty Regulation, availability, compute access and sovereignty	36-38
08	Watchlist, Decision Triggers and Previous Signal Review Forward triggers and review of previous signals	39
09	Methodology, Source Register and Corrections Definitions, calculations, source register and correction record	40-41
10	Publication back matter Publication information and contact	42

How to read this report.

The report separates evidence, interpretation and decision relevance.
Read in that order, then test the watch condition in the next edition.

READ IN THIS ORDER

- 01

Finding

The directly observed, reported or calculated result.
- 02

Assessment

yellow3's interpretation of what changed and why it matters.
- 03

Decision implication

The question the evidence creates for an organization or policymaker.
- 04

Watch condition

A measurable trigger that will be checked again in the next edition.

NO MATERIAL CHANGE

A valid published outcome. The monitored page remains in place when conditions are stable.

READING RULES

- 01

Color identifies categories, not good or bad.
- 02

Rank movement shows direction, not recommendation.
- 03

Every material figure carries a local source.
- 04

pp means percentage points; n= states sample size.

METHODOLOGY STATUS

VERSION	MI-MET-v1.0
SERIES	STANDARD SERIES
CHANGE THIS EDITION	NONE RECORDED
HISTORICAL REVISION	NONE
DATA CUT	13 Jul 2026, 10:42 UTC

Full definitions, calculations and source register appear on pages 40-41.

The market moved again.

Asia extended its lead, US share fell, and a 4% output-price cut added a new economic signal to the top tier.

ASIA SHARE	US SHARE	NUMBER ONE	BIGGEST RANK MOVE
60.60%	32.67%	Tencent Hunyuan	NVIDIA Nemotron
+1.03 pp	-2.73 pp	(Pro) / 2-week streak	+6 ranks

THE DECISION BRIEF

FINDING / ASSESSMENT / DECISION / WATCH

01 Regional power shift

FINDING / OBSERVED / MI.ADOPTION.REGIONAL_SHARE

Asia reached 60.60%, up 1.03 pp. US share fell to 32.67%; Europe remained at 0.57%.

ASSESSMENT

Asia's advantage now affects procurement assumptions.

DECISION IMPLICATION

Benchmark Asian and US leaders within the same procurement set.

WATCH CONDITION

Asia remains above 60% at the next data cut.

02 Leadership remains fluid

FINDING / OBSERVED / MI.MODEL.RANK

Tencent Hunyuan (Pro) held number one for two weeks. NVIDIA Nemotron rose six ranks.

ASSESSMENT

The leaderboard is still rotating quickly.

DECISION IMPLICATION

Refresh model shortlists monthly rather than relying on annual rankings.

WATCH CONDITION

The number-one model holds four consecutive weeks.

03 Price enters the adoption contest

FINDING / REPORTED / MI.PRICE.CHANGE

DeepSeek V3 Flash output pricing fell 4% during the reporting period.

ASSESSMENT

Price is becoming an active adoption lever.

DECISION IMPLICATION

Compare standardized workload cost alongside capability, reliability and access.

WATCH CONDITION

A second top-five model cuts output pricing.


Sources: yellow3 live adoption dataset and provider pricing records. Data cut: 13 Jul 2026, 10:42 UTC.

A regional lead without a dominant model.

Asian-built models occupy every top-five position, but no single model controls the field. Only 0.36 percentage points separate first and second.

TOP FIVE MODELS BY OBSERVED ROUTED SHARE

OBSERVED / MI.MODEL.TOP5_SHARE



CONCENTRATION SNAPSHOT

TOP FIVE COMBINED

47.05%

TOP TWO GAP

0.36 pp

ASIA-BUILT IN TOP FIVE

5 / 5

MOVED THIS WEEK

4 / 5

ASSESSMENT

Regional concentration is high; individual model concentration is not. The market can change leaders without changing its regional centre of gravity.

DECISION IMPLICATION

Do not treat the leading region as one interchangeable supplier. Maintain independent model and provider exit paths.

WATCH CONDITION A model exceeds 15% routed share or holds number one for four consecutive weeks.

Source: yellow3 live adoption dataset. Data cut: 13 Jul 2026, 10:42 UTC.

The same evidence creates three decisions.

Regional concentration, close model competition and price movement do not mean the same thing to every reader.

ASIA SHARE

60.60%

ASIA-BUILT IN TOP FIVE

5 / 5

TOP TWO GAP

0.36 pp

EUROPE SHARE

0.57%

DECISION LENS

INFERRED / CONFIDENCE: MEDIUM

01 BUSINESS

ASSESSMENT

Model choice is moving faster than annual procurement cycles.

DECISION IMPLICATION

Maintain a multi-provider shortlist and require an exit path for every production model.

WATCH CONDITION

A leader holds four weeks or one provider exceeds 15%.

PRIMARY EXPOSURE

PROCUREMENT

02 JOURNALISM

ASSESSMENT

Regional dominance is not the same as single-model dominance.

DECISION IMPLICATION

Report region, provider and model concentration as separate measures.

WATCH CONDITION

A model exceeds 15% or the top-five regional mix changes.

PRIMARY EXPOSURE

PUBLIC NARRATIVE

03 POLICY

ASSESSMENT

Europe's weakness is adoption and distribution, not only model supply.

DECISION IMPLICATION

Track access, routing and enterprise deployment beside model releases.

WATCH CONDITION

Europe crosses 1% routed share or enters the global top ten.

PRIMARY EXPOSURE

STRATEGIC AUTONOMY

Evidence base: pages 5-6 and yellow3 live adoption dataset. Assessment class: INFERRED.
Data cut: 13 Jul 2026, 10:42 UTC.

Act on the shift, not on a permanent winner.

The July evidence supports changes in procurement, evaluation and reporting.
It does not support locking into one model.

DECISION REGISTER / JULY 2026

INFERRED / CONFIDENCE SHOWN PER DECISION

DECISION / LENS / HORIZON / CONFIDENCE		EVIDENCE	ACTION
01	Refresh the model shortlist monthly. BUSINESS / NOW / HIGH	0.36 pp separates #1 and #2; four of the top five moved.	Re-run the approved-model evaluation before each procurement cycle.
02	Require an exit path for every production model. BUSINESS / 30 DAYS / MEDIUM	Leadership rotates while regional concentration remains high.	Document fallback model, data export and prompt compatibility.
03	Standardize workload cost. BUSINESS / 30 DAYS / MEDIUM	A reported 4% output-price cut changed the cost signal.	Compare monthly workload cost, not token list price alone.
04	Separate three kinds of concentration. JOURNALISM / NEXT REPORT / HIGH	Asia holds 60.60%; the leading model holds 11.66%.	Report region, provider and model concentration separately.
05	Measure European adoption, not release count. POLICY / QUARTERLY / MEDIUM	Europe holds 0.57% and no top-five position.	Add routed share, access and deployment to the policy scorecard.

NEXT FORMAL REVIEW	01 Aug 2026	ACCOUNTABILITY	Re-score against the watch conditions.
--------------------	-------------	----------------	--


Evidence base: pages 5-7. Data cut: 13 Jul 2026, 10:42 UTC.

Asia's usage exceeds its model footprint.

Half of the top twenty models are Asian-built, yet they capture 60.60% of routed usage.
US models show the inverse pattern: 45% of representation, but 32.67% of usage.

USAGE CAPTURE VERSUS TOP-TWENTY REPRESENTATION

CALCULATED FROM OBSERVED / PROVISIONAL



CAPTURE PREMIUM AND DEFICIT

CALCULATED / MI.ADOPTION.REPRESENTATION_VARIANCE

ASIA CAPTURE PREMIUM	US CAPTURE DEFICIT	ASIA / US YIELD RATIO
+10.60 pp	-12.33 pp	1.67x
60.60% usage vs 50% of models	32.67% usage vs 45% of models	6.06% vs 3.63% per listed model

ASSESSMENT

Representation does not explain Asia's usage lead.
Asian-built models over-index on routed usage while
US-built models materially under-index.

DECISION IMPLICATION

Treat model count as footprint, not market power.
Use capture variance to identify regions converting
representation into actual demand.

WATCH CONDITION Asia's premium falls below 5 pp or the US deficit narrows to less than 5 pp.


Source: yellow3 live adoption dataset. Benchmark = regional model count / 20. Data cut: 13 Jul 2026, 10:42 UTC.

Two-thirds of routed usage sits in the top ten.

The first five models account for 47.05%; ranks six to ten add another 19.63%.
One in five routed tokens still falls outside the reported top twenty.

CUMULATIVE ROUTED SHARE BY REPORTED RANK

CALCULATED / MI.ADOPTION.CUMULATIVE_SHARE



CONCENTRATION BREAKPOINTS

CALCULATED / REPORTED TOP-TWENTY SHARE

TOP FIVE	RANKS 6-10	RANKS 11-20	OUTSIDE TOP 20
47.05%	+19.63 pp	+13.41 pp	19.91%
9.41% average per model	3.93% average per model	1.34% average per model	not individually listed

ASSESSMENT

Usage is concentrated, but the curve does not close around a small elite. Nearly one-third sits beyond the top ten and one-fifth beyond the top twenty.

WATCH CONDITION Top-ten share exceeds 75% or top-twenty share exceeds 90%.

Source: yellow3 live adoption dataset. Cumulative values calculated from reported shares. Data cut: 13 Jul 2026, 10:42 UTC.

DECISION IMPLICATION


Cover the leaders and preserve a challenger lane.
A top-five-only shortlist misses a material share of observed demand.

Most rank changes were local. One was not.

Sixteen of eighteen continuing top-twenty models moved by one place or less.
NVIDIA Nemotron climbed from #14 to #8; two models entered the ranking.

DISTRIBUTION OF SEVEN-DAY RANK MOVEMENT

OBSERVED / PROVISIONAL / MI.MODEL.RANK_MOVEMENT



MOVEMENT COMPOSITION

OBSERVED / REPORTED TOP TWENTY

GAINED RANKS	UNCHANGED	LOST RANKS	NEW ENTRANTS
6	5	7	2
one moved six places	including the leader	all fell one place	entered at #18 and #20

ASSESSMENT

The ranking is active, not chaotic. Most changes are local; Nemotron's six-place rise is the clear exception requiring explanation.

WATCH CONDITION Another model moves three or more ranks or top-three membership changes.

Source: yellow3 live adoption dataset. Movement compares reported weekly ranks. Data cut: 13 Jul 2026, 10:42 UTC.

DECISION IMPLICATION


Set review thresholds around material movement.
Investigate moves of three or more ranks and new entrants reaching the top ten.

Asia is gaining at scale. The US is not.

Asia holds 60.60% and rose 1.03 pp; the US holds 32.67% and fell 2.73 pp.
Other gained from a small base; Europe remained below 1% and edged lower.

REGIONAL SCALE-MOMENTUM MATRIX

OBSERVED / PROVISIONAL / MI.ADOPTION.SHARE_AND_CHANGE



NOTICE: Reported regional deltas sum to -0.64 pp; provisional values may reflect rounding, timing and category updates.

SCALE AND MOMENTUM READOUT

OBSERVED / CALCULATED

MATERIAL-SCALE REGIONS	REGIONS GAINING	ASIA-US MOMENTUM GAP	REPORTED DELTA SUM
2	2	3.76 pp	-0.64 pp
Asia and United States	Asia and Other	+1.03 minus -2.73 pp	requires reconciliation

ASSESSMENT

Asia and the US are the only regions at material scale, but their weekly directions diverged. Other is gaining without scale; Europe has neither.

WATCH CONDITION Other crosses 10% routed share or the US weekly decline narrows below 1 pp.

Source: yellow3 live adoption dataset. Material-scale threshold: 10%. Data cut: 13 Jul 2026, 10:42 UTC.

DECISION IMPLICATION

Separate scale from momentum in regional reporting.
Do not infer exact token transfers until the reported regional deltas reconcile.

Rank direction is not the same as demand.


Models that slipped one rank still hold 32.99% of all routed usage, the largest cohort.
Upward movers hold 20.64%; new entrants account for 1.89%.

ROUTED-USAGE PARTITION BY WEEKLY RANK STATUS

CALCULATED FROM OBSERVED / PROVISIONAL

CIRCLE AREA PROPORTIONAL TO ROUTED SHARE

TOTAL ROUTED SHARE = 100%



NOTICE: Rank status is ordinal. A model can fall in rank without losing routed share.

Cohorts use current share and reported 7D rank status.

DEMAND BEHIND MOVEMENT

CALCULATED / MI.ADOPTION.SHARE_BY_RANK_STATUS

CURRENT TOP-20 SHARE	CONTINUING MODEL SHARE	MOVED MODEL SHARE	DOWN-UP SHARE GAP
80.09%	78.20%	53.63%	12.35 pp
includes two new entrants	18 ranked models	gained or lost rank	32.99% minus 20.64%

ASSESSMENT

Ordinal decline does not equal demand contraction.
The seven models that fell still represent the largest single share cohort in the current market.

DECISION IMPLICATION

Track rank change and share change separately.
Escalate a falling model only when routed share also contracts or the rank decline persists.

WATCH CONDITION A model falls for three consecutive cuts or loses more than 2 pp routed share.

Source: yellow3 live adoption dataset. Cohort totals calculated from reported top-twenty shares. Data cut: 13 Jul 2026, 10:42 UTC.

Visible momentum is not the same as adoption.

Evidence strengthens as it moves from attention to repeated production use.
yellow3 routed-token data belongs at Level 03: Usage, bounded to its contributing source.

ADOPTION EVIDENCE LADDER

INFERRED / CONFIDENCE: HIGH

LEVEL / STRENGTH	EVIDENCE	CAN SUPPORT	DOES NOT PROVE
05 PRODUCTION HIGH	Repeated workloads, retention, workflow integration, outcomes	Sustained operational use in the observed workload	Fit beyond the observed organization or workload
04 PROCUREMENT HIGH / DOCUMENTED	Contracts, spend, paid seats, approved shortlists	Commercial commitment and buying intent	Repeated production use or realized business value
03 USAGE MEDIUM / SOURCE-BOUNDED	API tokens, active users, workload share, retention	Observed activity inside the declared source	Global market share, contract, production use or value
YELLOW3 ROUTED-TOKEN DATA SITS HERE		SOURCE-SPECIFIC USAGE, NOT GLOBAL MARKET SHARE	
02 EXPERIMENTATION LOW-MEDIUM	Playground use, downloads, GitHub activity, trials	Testing and developer interest	Paid use, retention or deployment
01 ATTENTION LOW	Searches, web traffic, media and social activity	Visibility and public salience	Evaluation, usage or adoption

DATASET BOUNDARY / CURRENT EDITION

MI-DATA-2026-07-13

UNIVERSE	PERMITTED CLAIM	PROHIBITED CLAIM
Contributing routed-token dataset at the declared snapshot	Observed routed share within this dataset	Total global model market share or enterprise adoption

ASSESSMENT

Routed tokens are stronger evidence than attention or trials, but remain bounded usage evidence.

WATCH CONDITION A Level 03 usage signal is described as adoption without Level 04 or 05 corroboration.

Sources: [1] OpenRouter, State of AI; [2] GitHub, Starring; [3] Hugging Face Hub, Download Statistics; [4] yellow3 MI-DATA-2026-07-13.

DECISION IMPLICATION

Use routed share to monitor direction and concentration.
Require procurement or production evidence before claiming enterprise adoption.

The ranking contains one identity collision.

Two records use the reported label Gemini 2.5 Flash Lite but carry different shares and movements. Both records remain separate. Endpoint equivalence is not assumed.

MODEL MARKET LEDGER / REPORTED TOP TWENTY

OBSERVED / SOURCE-BOUNDED USAGE

RANKED RECORDS	UNIQUE LABELS	TOP-TWENTY SHARE	IDENTITY REVIEW
20	19	80.09%	2
current source rows	one repeated label	of this dataset	records unresolved

#	Y3 RECORD	PROVIDER	REPORTED MODEL LABEL	SHARE	7D	LABEL STATE
1	R001	Tencent	Hunyuan (Pro)	11.66%	0	UNIQUE
2	R002	Xiaomi	MiMo V2.5	11.30%	+1	UNIQUE
3	R003	DeepSeek	V3 Flash	9.93%	-1	UNIQUE
4	R004	MiniMax	M2	8.09%	-1	UNIQUE
5	R005	Z.ai	GLM 5.2	6.07%	-1	UNIQUE
6	R006	DeepSeek	V3 Pro	5.37%	0	UNIQUE
7	R007	Anthropic	Claude 4.7 Opus	4.29%	0	UNIQUE
8	R008	NVIDIA	Nemotron 3 Ultra 8B (free)	3.80%	+6	UNIQUE
9	R009	Anthropic	Claude 4.8 Opus	3.86%	-1	UNIQUE
10	R010	StepFun	Step 3.7 Flash	2.31%	-1	UNIQUE
11	R011	Anthropic	Claude 4.8 Sonnet	1.89%	-1	UNIQUE
12	R012	Google	Gemini 2.5 Flash Preview	1.70%	+1	UNIQUE
13	R013	Anthropic	Claude Sonnet 3.5	1.75%	+1	UNIQUE
14	R014	Qwen	Qwen2.5 GPT 5.5	1.66%	-1	UNIQUE
15	R015	Google	Gemini 2.5 Flash	1.36%	0	UNIQUE
16	R016	Xiaomi	MiMo V2.5 Pro	1.13%	+2	UNIQUE
17	R017	Google	Gemini 2.5 Flash Lite*	1.07%	-1	REPEATED
18	R018	Psalado	Laguna 8B (Instruct)	1.03%	NEW	UNIQUE
19	R019	Google	Gemini 2.5 Flash Lite*	0.96%	+1	REPEATED
20	R020	Qwen	Qwen2.5 GPT Q6S 128G	0.86%	NEW	UNIQUE

* Repeated reported label. R017 and R019 remain separate until endpoint, host and dated-version metadata resolve equivalence. Reported order retained; shares are not re-ranked.

DATASET BOUNDARY

Source-specific routed-token usage at the declared cut.
It is not total global market share or enterprise adoption.

IDENTITY CONTROL

Present: provider, reported label, rank, share and data cut.
Absent: endpoint, host, region, runtime and lifecycle status.

DECISION IMPLICATION Monitor position here. Resolve endpoint identity before capability, cost, risk or procurement comparison.

WATCH CONDITION A repeated alias changes independently, or an undated endpoint changes without a new label.

Source: yellow3 MI-DATA-2026-07-13; contributing routed-token source labels. Data cut: 13 Jul 2026, 10:42 UTC.

Two models entered. No release was proven.

NEW means first appearance in the source ranking, not a launch date, GA status or availability change. The current routed source supplies no authoritative deprecation or retirement record.

LIFECYCLE SIGNAL REGISTER / CURRENT CUT

OBSERVED + REPORTED / PROVISIONAL

SIGNAL / STRENGTH	SOURCE RECORD	WHAT IS OBSERVED	WHAT IS NOT ESTABLISHED
01 USAGE ENTRY MEDIUM SOURCE-BOUNDED	R018 / PSALADO Laguna 8B (Instruct) #18 / 1.03% / NEW	First appearance in the current top twenty	Release date, GA, endpoint, licence or production access
02 USAGE ENTRY MEDIUM SOURCE-BOUNDED	R020 / QWEN Qwen2.5 GPT Q6S 128G #20 / 0.86% / NEW	First appearance in the current top twenty	Release date, GA, endpoint, licence or production access
03 LABEL QUALIFIER LOW LABEL-ONLY	R012 / GOOGLE Gemini 2.5 Flash Preview #12 / 1.70% / +1	The source label contains PREVIEW	Current lifecycle state, GA path or retirement date
04 ACCESS QUALIFIER LOW LABEL-ONLY	R008 / NVIDIA Nemotron 3 Ultra 8B (free) #8 / 3.80% / +6	The source label contains FREE	Price, duration, service tier or production terms

Register classifies signals carried by the routed dataset. It does not infer missing provider lifecycle events.

Cut: 13 Jul 2026, 10:42 UTC

LIFECYCLE COVERAGE / CURRENT SOURCE

NOT SUPPLIED IS NOT EVIDENCE OF NO EVENT

RELEASE / GA DATE	UPGRADE / ALIAS EVENT	DEPRECATION DATE	RETIREMENT DATE
NOT SUPPLIED	NOT SUPPLIED	NOT SUPPLIED	NOT SUPPLIED


ASSESSMENT	DECISION IMPLICATION
This source carries ranking status, not a lifecycle log. The two entrants are valid usage signals, but the release, availability and retirement record remains incomplete. WATCH CONDITION A ranked model moves to GA, is deprecated, retires or changes access within 30 days. Source: yellow3 MI-DATA-2026-07-13. Lifecycle fields assessed against the Y3-RS-v1.0 model identity standard.	Use entrants as monitor and test triggers. Require official status, endpoint and contract evidence before migration, procurement or retirement decisions.

Two close scores. No capability delta.

MiMo V2.5 and V2.5 Pro are different variants, and their public ClawEval measures are described differently. No movement is attributed until model, task, harness, settings and scoring are comparable.

PUBLIC CAPABILITY SIGNAL / MIMO FAMILY

REPORTED / PROVIDER-SUPPLIED



CAPABILITY MOVEMENT ELIGIBILITY GATE

ALL FIVE CONDITIONS REQUIRED

REQUIREMENT		CURRENT EVIDENCE	STATUS
01	Exact model and version match	V2.5 versus V2.5 Pro	FAIL
02	Same task and subset	General subset versus Pass^3	NOT ESTABLISHED
03	Same harness and tool access	Common configuration not disclosed	NOT ESTABLISHED
04	Same scoring and trial count	Definitions and n not aligned	NOT ESTABLISHED
05	Independent replication	Provider records only	NO

CAPABILITY MOVEMENT STATUS

WITHHELD

No month-on-month delta calculated

CAN SUPPORT	DOES NOT SUPPORT
Two provider-published agentic capability signals. A justified trigger for matched internal testing.	A 1.7-point gain, cross-variant superiority or improvement on yellow3 business workloads.
DECISION IMPLICATION Run both variants through one frozen workload harness before selection or migration.	

Sources: Xiaomi MiMo-V2.5 release note, 22 Apr 2026; Xiaomi MiMo-V2.5-Pro release note, 27 Apr 2026.

Evidence: REPORTED

Uncertainty makes a benchmark decision-grade.

MiMo V2.5 Pro publishes a model tag, benchmark name and approximate token budget.
Trial count, confidence interval, failure categories and independent replication are absent from the cited record.

BENCHMARK DISCLOSURE AUDIT / MIMO V2.5 PRO

REPORTED INPUT / INFERRED CONFIDENCE: LOW

01 EXACT MODEL PRESENT mimo-v2.5-pro dated 27 Apr 2026	02 TASK DEFINITION PARTIAL ClawEval named task set not listed	03 HARNESS / TOOLS PARTIAL agentic context stated config not reproducible	04 TOKEN BUDGET PRESENT approximately 70K tokens / trajectory
05 TRIAL COUNT ABSENT number of runs not stated	06 VARIANCE / CI ABSENT no interval or dispersion reported	07 FAILURE CATEGORIES ABSENT no error or failure taxonomy	08 INDEPENDENT CHECK ABSENT provider record only

DISCLOSURE RESULT

CONTROL COMPLETENESS / NOT A PERFORMANCE SCORE

PRESENT 2 fully disclosed	PARTIAL 2 some evidence	ABSENT 4 not in source	DECISION CONFIDENCE LOW yellow3 assessment
--	--	---	---

BENCHMARK USE PERMISSION

CURRENT SOURCE / CURRENT CUT

DECISION USE	PERMISSION	EVIDENCE CONSEQUENCE
Candidate discovery	PERMITTED	sufficient to enter the test queue
Internal test prioritization	PERMITTED	use as a hypothesis, not a result
Procurement or migration	WITHHELD	uncertainty and replication are missing

ASSESSMENT Informative signal; uncertainty-aware comparison is not yet possible.

DECISION IMPLICATION Require repeated matched trials, variance and failure categories before ranking or procurement.

Sources: Xiaomi MiMo-V2.5-Pro release note, 27 Apr 2026; yellow3 benchmark control standard. Retrieved 14 Jul 2026.

A score shows how much. It does not show how it failed.

Failure categories are absent from the cited provider benchmark record.
Six operationally different failures therefore remain indistinguishable behind the aggregate.

BENCHMARK FAILURE-MODE REGISTER

REPORTED SOURCE / FAILURE DATA ABSENT

MODE	FAILURE CATEGORY	EVIDENCE IN SOURCE	DECISION CONSEQUENCE
01	Incorrect result	Not supplied	Accuracy risk cannot be isolated
02	Incomplete task	Not supplied	Partial completion can pass unnoticed
03	Tool misuse	Not supplied	Agent permission risk remains unknown
04	Citation failure	Not supplied	Verification burden cannot be priced
05	Timeout or system error	Not supplied	Production reliability is unquantified
06	Policy refusal	Not supplied	Workload coverage cannot be predicted

CATEGORIES	SUPPLIED	AGGREGATE SCORE	FAILURE CONFIDENCE
6	0	64%	LOW
required taxonomy	in cited record	provider reported	yellow3 assessment

ASSESSMENT

An aggregate benchmark score cannot reveal whether failure is factual, operational, safety-related or infrastructural.

DECISION IMPLICATION

Require a failure taxonomy and example traces before translating the score into controls, staffing or service levels.

WATCH CONDITION

A repeated matched test publishes failure counts by category.

Source: Xiaomi MiMo-V2.5-Pro release note, 27 Apr 2026; yellow3 benchmark control standard.

Four signals changed testing. None changed procurement.

Identity, lifecycle and benchmark evidence all produced actions this month.
The actions are verification and matched testing, not model replacement.

MODEL DECISION DELTA / JULY 2026

INFERRED / CONFIDENCE: HIGH

DELTA	NEW EVIDENCE	DECISION AFFECTED	RECOMMENDED RESPONSE
01	Repeated Gemini Flash Lite label	Model comparison	Resolve endpoints before merging records
02	Two new ranking entrants	Research coverage	Add both to the monitoring queue
03	MiMo public scores not comparable	Capability shortlist	Run one frozen workload harness
04	Four benchmark controls absent	Procurement confidence	Withhold migration recommendation

SIGNALS REVIEWED	TEST ACTIONS	PROCUREMENT CHANGES	MODEL SWITCHES
4 pages 15-19	3 verify or measure	0 this cut	0 recommended

ASSESSMENT

Evidence changed the research sequence but did not clear the threshold for a purchasing or migration decision.

DECISION IMPLICATION

Keep current deployments stable. Resolve identity and run matched tests before changing a shortlist or contract.

WATCH CONDITION Any signal gains authoritative lifecycle, matched benchmark or production evidence.


Sources: yellow3 pages 15-19; assessment frozen at the 13 Jul 2026 data cut.

No provider exceeds 15.30% of routed usage.

DeepSeek leads the provider view; Xiaomi, Anthropic and Tencent follow within 3.64 percentage points. The top five known providers account for 59.27% of this source-specific dataset.

KNOWN PROVIDER SHARE / REPORTED TOP TWENTY

CALCULATED FROM OBSERVED / SOURCE-BOUNDED



TOP PROVIDER	TOP THREE	TOP FIVE	KNOWN PROVIDERS
15.30%	39.52%	59.27%	11
DeepSeek	known providers	known providers	top-twenty records

ASSESSMENT	DECISION IMPLICATION
Provider concentration is material but distributed. The regional lead does not collapse into one supplier.	Maintain at least two provider-qualified alternatives; concentration alone does not justify a single-provider strategy.


WATCH CONDITION A known provider exceeds 20% source-bounded routed share.

Source: yellow3 MI-DATA-2026-07-13. Provider totals calculated from reported top-twenty shares.

Portfolio breadth is not the same as provider strength.

Anthropic and Google each place four records in the top twenty.
DeepSeek leads known provider share with two records; Tencent reaches fourth with one.

PROVIDER PORTFOLIO BREADTH / SHARECALCULATED FROM OBSERVED / SOURCE-BOUNDED



WIDEST PORTFOLIOS	SHARE LEADER	SINGLE-MODEL SHARE	OUTSIDE TOP 20
4	15.30%	11.66%	19.91%
Anthropic / Google	DeepSeek	Tencent	provider unknown

ASSESSMENT

A broad portfolio can diversify use cases; a concentrated portfolio can still command significant demand.

DECISION IMPLICATION

Assess provider dependence at the model and portfolio level. Do not use record count as a proxy for durability.

WATCH CONDITION One provider combines more than 20% routed share with four or more ranked records.

Source: yellow3 MI-DATA-2026-07-13. Model counts and provider shares calculated from reported rows.

Xiaomi is pairing model breadth with access optionality.

The provider released two V2.5 variants, exposed first-party API access and published open weights. A later UltraSpeed route traded a reported threefold price for approximately tenfold output speed.

XIAOMI STRATEGY SIGNALS / APR-JUN 2026

REPORTED / PROVIDER-SUPPLIED

DATE	SIGNAL	EVIDENCE	STRATEGIC READING
22 APR	MiMo V2.5 released	API, 1M context, open weights	Broad multimodal access
27 APR	MiMo V2.5 Pro released	API and open-source rollout	Premium agentic tier
27 APR	No pricing change reported	Existing surfaces retained	Upgrade without list-price reset
08 JUN	UltraSpeed access route	Limited access; 3x price, ~10x speed	Latency segmentation

MODEL VARIANTS	ACCESS SURFACES	CONTEXT	STRATEGY
2 V2.5 / V2.5 Pro	3 API / studio / weights	1M provider reported	OPTIONALITY cost / speed / control

ASSESSMENT

The strategy expands choice across capability, speed and deployment control. All claims remain provider-reported.

DECISION IMPLICATION

Test whether access optionality reduces switching risk; verify pricing, licence and service terms before contracting.

WATCH CONDITION A first-party access term or open-weight licence changes.

Sources: Xiaomi MiMo-V2.5 and V2.5-Pro release notes; Xiaomi UltraSpeed note, 8 Jun 2026.

No material partnership change was verified at this cut.

Product and access signals are present, but no authoritative new contract or partnership record is. The page remains in place so absence is visible and comparable next month.

PARTNERSHIP AND COMMERCIAL SIGNAL REGISTER

NO MATERIAL CHANGE / FINAL FOR THIS CUT

NO MATERIAL CHANGE

No verified M&A, strategic partnership, enterprise contract, capital event or infrastructure commitment entered the evidence set.

MONITORED FIELD	CURRENT STATUS	REQUIRED EVIDENCE
Strategic partnership	No verified record	Official agreement or filing
Enterprise win	No verified record	Named contract or customer record
Capital / M&A	No verified record	Filing, audited record or parties
Infrastructure commitment	No verified record	Capacity, location, date and counterparty

ASSESSMENT

A blank commercial register is a valid outcome. It must not be filled with rumour or weak social signals.

DECISION IMPLICATION

Maintain current provider-durability assumptions until a primary commercial record changes them.

WATCH CONDITION A primary source confirms a contract, partnership, capital or infrastructure event.

Source: yellow3 evidence foundation reviewed through 14 Jul 2026. No qualifying record identified.

Usage visibility is stronger than provider-durability evidence.

Five leading providers are visible in the routed dataset.
Only Xiaomi has lifecycle and deployment-control records in the current evidence foundation.

PROVIDER DURABILITY / EVIDENCE COMPLETENESS

OBSERVED + REPORTED / NOT A DURABILITY RANKING

	ROUTED PRESENCE	PORTFOLIO >1	LIFECYCLE RECORD	OPEN / SELF-HOST	ENTERPRISE TERMS
DeepSeek	YES	YES	NOT IN SET	NOT IN SET	NOT IN SET
Xiaomi	YES	YES	PRESENT	PRESENT	NOT IN SET
Anthropic	YES	YES	NOT IN SET	NOT IN SET	NOT IN SET
Tencent	YES	NO	NOT IN SET	NOT IN SET	NOT IN SET
MiniMax	YES	NO	PARTIAL	PRESENT	NOT IN SET

PROVIDERS	USAGE VISIBLE	TERMS COMPLETE	DURABILITY RANK
5	5 / 5	0 / 5	WITHHELD
leading known share	source-bounded	current foundation	evidence incomplete

ASSESSMENT	DECISION IMPLICATION
Observed usage does not establish provider financial, contractual or operational durability.	Use the matrix to target diligence. Do not rank provider durability until terms, lifecycle and commercial records are complete.

WATCH CONDITION A provider supplies audited financials, enterprise terms and lifecycle commitments.

Sources: yellow3 MI-DATA-2026-07-13; Xiaomi and MiniMax primary model records. Evidence-set status only.

Token price is one line in a twelve-line cost ledger.

The routed dataset contains no normalized provider-pricing fields.
A comparable economics series is therefore withheld until every material charge is captured.

PRICING COMPONENT COVERAGE / CURRENT FOUNDATION

OBSERVED ABSENCE / FINAL FOR THIS CUT

01	02	03	04
STANDARD INPUT	STANDARD OUTPUT	CACHE READ	CACHE WRITE
NOT NORMALIZED	NOT NORMALIZED	NOT NORMALIZED	NOT NORMALIZED
05	06	07	08
BATCH DISCOUNT	REASONING COMPUTE	LONG CONTEXT	TOOLS / SEARCH
NOT NORMALIZED	NOT NORMALIZED	NOT NORMALIZED	NOT NORMALIZED
09	10	11	12
STORAGE	CLOUD MARKUP	RESIDENCY PREMIUM	COMMITMENT / VOLUME
NOT NORMALIZED	NOT NORMALIZED	NOT NORMALIZED	NOT NORMALIZED

COMPONENTS REQUIRED	FULLY CAPTURED	PRICE COMPARISON	NEXT ACTION
12	0	WITHHELD	SOURCE
normalized fields	current foundation	not decision-grade	official rate cards

ASSESSMENT

A token-rate table would create false precision because caching, tools, residency and commitments are unobserved.

DECISION IMPLICATION

Build the complete ledger from dated official rate cards before comparing providers or forecasting spend.

WATCH CONDITION All twelve cost fields are captured for two or more comparable endpoints.

Sources: yellow3 MI-DATA-2026-07-13; pricing-field standard from MI-MET-v1.0.

Cost becomes comparable only when the work stays fixed.

yellow3 defines four recurring business workloads with stable inputs, outputs and pass criteria. The pack is specified this month; model measurements have not yet entered the series.

FIXED BUSINESS WORKLOAD PACK / MI-WORK-v0.1

DEFINED / NOT YET MEASURED

ID	WORKLOAD	FIXED INPUT AND OUTPUT	PASS CONDITION
W01	Board-pack memo	80-page pack; sourced executive memo	All critical issues cited and prioritized
W02	Contract comparison	Two supplier contracts; risk register	Material differences and obligations correct
W03	Spreadsheet analysis	Workbook; ten questions; checked output	Formulas, answers and audit trail pass
W04	Repository change	Issue, codebase, patch and test run	Required tests pass without regression

WORKLOADS	TRIALS	HARNESS	RESULTS
4	3	FROZEN	PENDING
recurring pack	per model / workload	MI-WORK-v0.1	no series yet

ASSESSMENT

A fixed workload converts model economics from token prices into repeatable business outcomes.

DECISION IMPLICATION

Freeze the pack and tool permissions before testing. Publish model, endpoint, location, trials and failures with every result.

WATCH CONDITION

The first two models complete all four workloads under the frozen harness.

Source: yellow3 workload standard MI-WORK-v0.1, defined 14 Jul 2026.

EUR1.00 per run becomes EUR2.24 per successful task.

The worked method includes reruns, tools and verification, then divides by successful outcomes. Figures are illustrative assumptions, not current provider pricing.

COST PER SUCCESSFUL TASK / WORKED METHOD

CALCULATED / ILLUSTRATIVE

BASE RUNS	100 x EUR1.00	EUR100
FAILED RERUNS	30 x EUR1.00	+ EUR30
TOOLS	fixed allowance	+ EUR20
VERIFICATION	human / automated	+ EUR40
TOTAL END-TO-END COST	100 attempted tasks	EUR190
SUCCESSFUL OUTCOMES	85 passed	/ 85

COST PER SUCCESSFUL TASK

EUR2.24

EUR190 total cost / 85 successful outcomes = EUR2.235, rounded to EUR2.24

ASSESSMENT

List-price cost per attempt understates the economic burden by 124% in this worked scenario.

DECISION IMPLICATION

Budget and compare models using cost per successful task. Keep every assumption visible and replace illustration with measured inputs.

WATCH CONDITION Measured retries, tools and verification change the multiplier by more than 10%.

Method: MI-ECON-v0.1. Illustrative assumptions only; not a provider or market estimate.

The higher-priced model is cheaper after quality adjustment.

A worked comparison combines run cost, verification burden and successful outcomes.
Model B costs twice as much per run but 11% less per successful task.

QUALITY-ADJUSTED COST / WORKED COMPARISON

CALCULATED / ILLUSTRATIVE

MODEL A	MODEL B
LOWER LIST PRICE	HIGHER LIST PRICE
RUN COST	RUN COST
EUR0.50	EUR1.00
SUCCESSFUL OUTCOMES	SUCCESSFUL OUTCOMES
60 / 100	90 / 100
VERIFICATION	VERIFICATION
EUR40	EUR20
COST / SUCCESS	COST / SUCCESS
EUR1.50	EUR1.33

LIST PRICE GAP	SUCCESS GAP	COST / SUCCESS	RESULT
+100%	+30	-11%	MODEL B
B versus A	B versus A	B versus A	lower effective cost

ASSESSMENT	DECISION IMPLICATION
Run price alone selects Model A; quality-adjusted cost selects Model B under the disclosed assumptions.	Require success and verification measurements before accepting a cheaper list-price model as the economic choice.

WATCH CONDITION

Measured success rates move either model's cost per success by more than 10%.

Method: MI-ECON-v0.1. Illustrative 100-task scenario; no provider identity attached.

A EURO.17 unit gap becomes EUR17,000 at scale.

The quality-adjusted example is projected across three standard monthly task volumes.
Volume amplifies small unit-economics differences; it does not improve evidence quality.

MONTHLY VOLUME SCENARIOS / QUALITY-ADJUSTED COST

CALCULATED / ILLUSTRATIVE



MODEL A / EUR1.50 PER SUCCESS

MODEL B / EUR1.33 PER SUCCESS

UNIT GAP	1K SAVING	10K SAVING	100K SAVING
EUR0.17	EUR170	EUR1,700	EUR17,000
per success	per month	per month	per month

ASSESSMENT	DECISION IMPLICATION
Small quality-adjusted differences become budget-relevant at operational scale.	Use three volume scenarios in every comparison and test sensitivity to success, verification and price changes.

WATCH CONDITION

Measured task volume or cost per success changes by more than 10%.

Method: MI-ECON-v0.1. Illustrative scenario based on page 29; not provider pricing.

The method is ready. The provider comparison is not.

Workloads, cost equations and volume scenarios are defined.
Official rate cards and measured success data remain outside the current foundation.

ECONOMICS DECISION GATE

INFERRED / CONFIDENCE: HIGH

GATE	REQUIRED EVIDENCE	CURRENT STATUS	CONSEQUENCE
01	Complete official pricing ledger	NOT CAPTURED	List-price comparison withheld
02	Frozen workload and tools	DEFINED	Testing can begin
03	Repeated successful outcomes	NOT MEASURED	Cost per success unavailable
04	Latency, retries and verification	NOT MEASURED	End-to-end cost unavailable
05	Volume and contract terms	NOT CAPTURED	Budget forecast unavailable

METHOD	WORKLOADS	MEASUREMENTS	ECONOMIC RANK
READY MI-ECON-v0.1	DEFINED MI-WORK-v0.1	PENDING no provider series	WITHHELD current cut

ASSESSMENT	DECISION IMPLICATION
A decision-grade economics comparison needs both market prices and observed workload outcomes.	Do not publish a cheapest-model ranking. Start the fixed workload series and archive dated official rate cards.

WATCH CONDITION

Two endpoints pass all five economics gates.

Sources: pages 26-30; MI-ECON-v0.1 and MI-WORK-v0.1. Assessment at 14 Jul 2026.

A capability score is not a production service level.

The cited benchmark record reports task performance and an approximate token budget. It does not provide latency distribution, error rate, uptime or regional performance.

PRODUCTION PERFORMANCE COVERAGE / MIMO V2.5 PRO

REPORTED SOURCE / SERVICE DATA ABSENT

01	02	03	04
TIME TO FIRST TOKEN	OUTPUT SPEED	COMPLETION TIME	P50 / P95 LATENCY
NOT SUPPLIED	NOT SUPPLIED	NOT SUPPLIED	NOT SUPPLIED
No comparable series	No comparable series	No comparable series	No comparable series

05	06	07	08
ERROR RATE	UPTIME	REGIONAL VARIANCE	RATE LIMITING
NOT SUPPLIED	NOT SUPPLIED	NOT SUPPLIED	NOT SUPPLIED
No comparable series	No comparable series	No comparable series	No comparable series

METRICS REQUIRED	SUPPLIED	SERVICE CONFIDENCE	PRODUCTION CLAIM
8	0	LOW	WITHHELD
production panel	cited record	yellow3 assessment	no service evidence

ASSESSMENT

Task capability may justify testing, but it cannot establish production speed, stability or capacity.

DECISION IMPLICATION

Collect endpoint-level p50 and p95 latency, errors, uptime and region data before setting service expectations.

WATCH CONDITION A provider or independent run supplies all eight metrics for 30 days.

Source: Xiaomi MiMo-V2.5-Pro release note, 27 Apr 2026; MI-PERF-v0.1 coverage standard.

No material reliability incident was verified at this cut.

The evidence foundation includes no authoritative incident feed for the ranked endpoints.
This is a coverage limitation, not evidence of uninterrupted service.

RELIABILITY AND INCIDENT REGISTER

NO MATERIAL CHANGE / COVERAGE LIMITED

NO VERIFIED INCIDENT

No qualifying outage, elevated-error period or breaking behaviour change entered the current evidence set.

MONITORED EVENT	CURRENT STATUS	REQUIRED PRIMARY RECORD
Outage	Not observable	Provider status incident with duration
Elevated error rate	Not observable	Timestamped endpoint metrics
Breaking behaviour	Not observable	Versioned change and reproducible test
Regional degradation	Not observable	Region-specific latency and error data

ASSESSMENT

A blank incident register is not a clean uptime record when status feeds are absent.

DECISION IMPLICATION

Connect authoritative status histories and endpoint telemetry before using this page for SLA or renewal decisions.

WATCH CONDITION A qualifying incident or thirty-day clean telemetry series enters the register.

Source: yellow3 evidence foundation through 14 Jul 2026. Endpoint status coverage not supplied.

MiMo V2.5 Pro is testable, not procurement-ready.

Technical access and open weights are documented in the provider record.
Privacy, residency, contract, administration and service evidence remain outside the foundation.

ENTERPRISE DEPLOYMENT READINESS / MIMO V2.5 PRO

REPORTED + INFERRED / CONFIDENCE: MEDIUM

CONTROL	REQUIRED EVIDENCE	CURRENT STATUS	READINESS EFFECT
01	Technical availability	PRESENT	Testing can begin
02	Open weights / licence	PRESENT	Self-host route available
03	Privacy / retention	NOT IN FOUNDATION	Data approval blocked
04	Residency / jurisdiction	NOT IN FOUNDATION	Regional approval blocked
05	DPA / contractual protection	NOT IN FOUNDATION	Procurement blocked
06	Admin / audit controls	NOT IN FOUNDATION	Operational control unknown
07	SLA / reliability	NOT IN FOUNDATION	Production commitment blocked
08	Lifecycle / deprecation	PARTIAL	Migration risk unresolved

CONTROLS	PRESENT	PARTIAL	READINESS
8	2	1	TESTABLE
deployment matrix	technical / weights	lifecycle	not procurement-ready

ASSESSMENT

Deployment access is established; enterprise protection and operational evidence are not.

WATCH CONDITION All eight readiness controls become present or contractually mitigated.

Sources: Xiaomi MiMo-V2.5-Pro release note; yellow3 EDR-v0.1 control matrix.

DECISION IMPLICATION


Permit controlled testing only. Hold procurement until privacy, residency, contract, admin, SLA and lifecycle controls are documented.

The evidence clears testing. Nothing beyond it.

Usage and provider capability claims justify a controlled evaluation.
Procurement and production remain blocked by economics, reliability and enterprise-control gaps.

ENTERPRISE FIT / EVIDENCE STAGE GATE

INFERRED / CONFIDENCE: HIGH



CURRENT POSITION

TEST

Four blocked gate exits remain

EXPLORE	TEST	PILOT	PROCURE + PROD
CLEARED	CURRENT	BLOCKED	BLOCKED
usage / identity	workload evidence	telemetry / privacy	terms / reliability

ASSESSMENT	DECISION IMPLICATION
The evidence supports learning, not commitment. Advancing stages without gate evidence would transfer uncertainty into operations.	Authorize a bounded test. Keep pilot, procurement and production approvals closed until their exit criteria are met.

WATCH CONDITION

Matched tests, economics and enterprise controls clear the pilot gate.

Sources: pages 15-34; stage logic defined in yellow3 EDR-v0.1.

Open weights create an exit route, not EU control.

Xiaomi and MiniMax document open-weight and first-party access paths in the evidence foundation. EU residency, governing terms and operational control remain unverified for every reviewed provider.

EUROPEAN CONTROL AND SOVEREIGNTY / EVIDENCE MATRIX REPORTED + OBSERVED / NOT A NATIONALITY RANKING

	ROUTED USE	FIRST-PARTY API	OPEN WEIGHTS	LICENCE	EU RESIDENCY	EXIT ROUTE
Xiaomi	OBSERVED	PRESENT	PRESENT	MIT	NOT STATED	PARTIAL
MiniMax	OBSERVED	PRESENT	PRESENT	MIT	NOT STATED	PARTIAL
DeepSeek	OBSERVED	NOT IN SET	NOT IN SET	NOT IN SET	NOT STATED	UNKNOWN
Anthropic	OBSERVED	NOT IN SET	NOT IN SET	NOT IN SET	NOT STATED	UNKNOWN
Google	OBSERVED	NOT IN SET	NOT IN SET	NOT IN SET	NOT STATED	UNKNOWN

PROVIDERS	OPEN-WEIGHT PATH	EU RESIDENCY	SOVEREIGNTY RANK
5	2	0 / 5	WITHHELD
reviewed	in foundation	verified	control incomplete

ASSESSMENT

Self-hostability improves switching control, but sovereignty also requires jurisdiction, data-location and operating evidence.

DECISION IMPLICATION

Treat open weights as one control dimension. Verify EU hosting, governing terms and exit feasibility before claiming sovereignty.

WATCH CONDITION A provider documents EU residency, governing jurisdiction and tested exit for one endpoint.


Sources: yellow3 routed dataset; Xiaomi and MiniMax primary model records; sovereignty matrix Y3-SOV-v0.1.

More buyer control means more operating responsibility.

Routed API, first-party API and self-hosted weights create different control and dependency profiles. No path is universally superior; each moves risk between provider, intermediary and buyer.

DEPLOYMENT CONTROL PATHS / DECISION MAP

INFERRED / CONFIDENCE: HIGH



BUYER CONTROL

INCREASES

ACCESS SPEED	CONTRACT CONTROL	OPERATOR CONTROL	BUYER BURDEN
HIGH	DIRECT	HIGH	HIGH
routed API	first-party API	self-hosted	self-hosted

ASSESSMENT	DECISION IMPLICATION
Deployment control is a trade: faster access delegates control; self-hosting returns control and operating burden to the buyer.	Select a path from data location, service, exit and operational requirements rather than model nationality.

WATCH CONDITION

A tested exit exercise shows a workload can move paths without material loss.

Source: yellow3 sovereignty framework Y3-SOV-v0.1; path characteristics are inferred.

Six evidence gaps now create six decision risks.

The risks arise from interpretation, identity, lifecycle, benchmarking, economics and reliability coverage. Severity is a yellow3 assessment; the underlying gaps are recorded elsewhere in this edition.

RECURRING RISK AND GOVERNANCE REGISTER

INFERRED / CONFIDENCE: HIGH

ID	RISK	LEVEL / FUNCTION	CURRENT CONTROL
R1	Source-bounded usage called global share	HIGH / GOVERN	Mandatory claim boundary on pages 14-15
R2	Repeated alias merged as one model	MEDIUM / MAP	Separate source records pending identity
R3	Lifecycle change missed	HIGH / MAP	Official status record required
R4	Incomparable benchmarks treated as delta	HIGH / MEASURE	Five-condition movement gate
R5	Token rate treated as total cost	HIGH / MEASURE	Twelve-field pricing ledger
R6	Capability treated as production reliability	HIGH / MANAGE	Production telemetry gate

OPEN RISKS	HIGH	MEDIUM	CLOSED
6	5	1	0
current edition	yellow3 assessment	identity collision	evidence still pending

ASSESSMENT

The dominant risk is not a single model failure; it is making decisions from incomplete or misclassified evidence.

WATCH CONDITION

A control closes, severity changes or a new material event enters the register.

Sources: pages 14-37; taxonomy aligned to NIST Govern, Map, Measure and Manage.

DECISION IMPLICATION

Keep all six controls active. Close a risk only when the missing source, measurement or contract evidence is archived.

Four decisions remain open at the next cut.

No evidence item from Pages 15-38 cleared a procurement or production threshold.
Each recommendation remains attached to a measurable reversal or confirmation condition.

HISTORICAL DECISION ACCOUNTABILITY / JULY BASELINE

INFERRED / CONFIDENCE: HIGH

ID	PRIOR DECISION	CURRENT STATUS	NEXT EVIDENCE / TRIGGER
D01	Keep repeated Gemini aliases separate	OPEN	Endpoint identity resolves R017 / R019
D02	Treat new entrants as usage signals	OPEN	Official release or GA record appears
D03	Withhold MiMo capability delta	OPEN	Matched independent benchmark is published
D04	Withhold procurement recommendation	OPEN	Economics, readiness and reliability clear
D05	Permit controlled testing	ACTIVE	Frozen workload pack produces first results

DECISIONS TRACKED	OPEN	ACTIVE	INVALIDATED
5	4	1	0
July baseline	evidence pending	controlled testing	no reversal yet

ASSESSMENT	DECISION IMPLICATION
This is the first accountability baseline. Future editions will show whether evidence confirmed, changed or invalidated each decision.	Carry every open decision into the August edition. Never delete a prior recommendation when the evidence changes.

WATCH CONDITION Any D01-D05 trigger is met before the August data cut.


Source: yellow3 decision register, Pages 15-38. Baseline established 14 Jul 2026.

The series is comparable because the method is frozen.

Universe, cutoff, evidence labels, calculations and decision gates are versioned with the edition.
A method change creates a new series or an explicit historical restatement.

METHOD FLOW / MI-MET-v1.0

FINAL / STANDARD SERIES



CONTROL	LOCKED VALUE	SERIES RULE
Universe / cut	MI-DATA-2026-07-13 / 13 Jul 10:42 UTC	Late data enters next edition
Evidence	Five-label evidence set / see Reading Guide	Confidence only on ESTIMATED / INFERRED
Benchmark	MI-WORK-v0.1 / three trials	Same endpoint, tools, settings and scoring
Economics	MI-ECON-v0.1	Total cost divided by successful outcomes
Identity	Provider, version, endpoint, region, status	Undated aliases never merged silently
Revision	STANDARD SERIES	Method change creates revision or break

ASSESSMENT

Comparability is a publication control, not a visual preference.
Every monthly edition keeps the same page roles and methods.

WATCH CONDITION MI-MET, MI-WORK, MI-ECON or identity rules change.

Methodology: MI-MET-v1.0. Standard series; no methodology change recorded this edition.

DECISION IMPLICATION

Use this edition as the July baseline. Flag any future method change before comparing movement.

Every recurring source needs evidence and reuse rights.

Primary records support the strongest claims; commercial reuse terms remain source-specific. Corrections are versioned and permanent. None are recorded for this edition.

SOURCE AND RIGHTS REGISTER / CURRENT EDITION

FINAL REGISTER

ID	SOURCE GROUP	EVIDENCE ROLE	RIGHTS / CONTROL STATUS
S01	yellow3 MI-DATA-2026-07-13	Routed usage snapshot	Internal source register
S02	Xiaomi official model records	Release, access and claims	Public primary / linked
S03	MiniMax GitHub and paper	Model, licence and benchmarks	Public primary / linked
S04	OpenRouter State of AI	Source-bounded usage context	Terms review required
S05	Independent benchmark methods	Comparability standard	Terms vary by source
S06	EU / OECD / NIST records	Policy and governance	Public authority / linked

SOURCE GROUPS	PRIMARY / AUTHORITY	RIGHTS REVIEW	CORRECTIONS
6	4	2	NONE
current register	strongest hierarchy	terms vary / review	recorded

ASSESSMENT	DECISION IMPLICATION
A source can be credible and still be unsuitable for automated commercial reuse. Evidence and rights are separate controls.	Complete rights review before recurring extraction. If any published value changes, issue a new version and correction record.
WATCH CONDITION	A source term changes or a correction is validated.
Correction record: none. Methodology change: none. Source-rights register updated 14 Jul 2026.	

The report ends. The decision record continues.

This edition is the July baseline for a fixed monthly intelligence series.
Data, alerts, exports and historical decisions remain attached to the publication identity.

PUBLICATION AND ACCESS RECORD

FINAL / Y3-MI-2026-07-v1.0

Y3-MI-2026-07-v1.0

yellow3 Model Intelligence Monthly Decision Report / July 2026

PUBLISHER	PUBLISHED	DATA CUT
yellow3 lab ApS	14 Jul 2026	13 Jul 2026, 10:42 UTC
REPORT STANDARD	METHODOLOGY	REPORT LENGTH
Y3-RS-v1.0	MI-MET-v1.0	42 pages

SUBSCRIBER RECORD

- Monthly report
- Data appendix
- Historical model records
- Pricing and change logs
- Critical alerts
- Exports and briefing notes

SUGGESTED CITATION	CORRECTIONS
yellow3 Model Intelligence Monthly Decision Report, Jul 2026, Y3-MI-2026-07-v1.0, page and figure number.	Never silent. Revisions receive a new version and permanent record.

Permanent record key: Y3-MI-2026-07-v1.0. Access through the yellow3 Model Intelligence subscriber archive.